

README — RL Trading Agent in Synthetic Multi-Good Market

This repository contains a **synthetic multi-good market environment** with **endogenous price dynamics** and a **reinforcement learning (PPO) trading agent**. Training logs are exported to Excel files and can then be analyzed to generate plots.

1) Project structure

The project must follow this structure:

```
PROJECT_ROOT/
├── game/                      # game / simulation code (Pygame)
│   ├── MainFrame.py            # main entry point for the game
│   ├── Solar_System.py
│   ├── objects.py
│   ├── camera.py
│   ├── drawing.py
│   └── UI_manager.py
    └── Images/                 # icons, UI images (play, pause, etc.)
└── ML/
    ├── ML_main.py              # main training launcher (runs multiple
    config)
    ├── prices_process.py       # price process parameters + market dynamics
    ├── ML_nd_priceprocess_functions.py # environment, logging, PPO helpers
    ├── Results_analysis.py     # data loading, aggregation, plots
    └── plots/                  # auto-generated figures
```

Important: all game-related scripts (`drawing.py`, `camera.py`, `UI_manager.py`, `Solar_System.py`, `objects.py`) **must be located in the same `game/` folder**. They rely on direct Python imports between each other.

2) Installation / requirements

This project has two parts: (i) the game/simulation (Pygame) and (ii) the RL + analysis stack.

External libraries (pip / conda)

Game (Pygame) - `pygame`

RL training + environment - `numpy` `pandas`
`filecrite turn0file1` `filecrite turn0file2` - `gymnasium` `stable-baselines3`

filecite turn0file1 filecite turn0file2 filecite turn0file0 - **torch** (PyTorch) filecite turn0file1
filecite turn0file2

Plots / analysis - **matplotlib** (incl. **matplotlib.patheffects**) filecite turn0file2
filecite turn0file3 - **openpyxl** (needed to read/write **.xlsx** logs via pandas)

Python standard library modules used

- **sys** filecite turn1file5
- **time** filecite turn0file1 filecite turn0file2
- **math** filecite turn1file4 filecite turn0file2
- **random** filecite turn1file4
- **copy** (**deepcopy**) filecite turn0file1
- **pathlib** (**Path**) filecite turn0file3
- **dataclasses** filecite turn0file2
- **typing** filecite turn0file2

3) Running the game (simulation)

The game / simulation is implemented using **Pygame** and serves as the interactive visual environment.

How to run the game

From the project root:

```
python game/MainFrame.py
```

This is the **only file you need to execute** to run the game.

- **MainFrame.py** initializes Pygame, the camera, UI, and the first solar system
- All other game scripts are imported automatically and must **not** be run individually

Game scripts overview

- **MainFrame.py** — main loop, event handling, rendering
- **Solar_System.py** — star systems, planets, entity management
- **objects.py** — planet and star classes
- **camera.py** — zooming, panning, world-to-screen conversion
- **drawing.py** — low-level drawing helpers
- **UI_manager.py** — UI elements, buttons, panels

If any of these files are moved or renamed, imports will break.

4) Training the models (RL)

Training is launched from **ML/ML_main.py**.

Run training

From the project root:

```
python ML/ML_main.py
```

This will train several models sequentially and produce Excel files such as:

```
model_value_lenght_60000_raw.xlsx  
model_value_lenght_40000_ent_coef_0.005_raw.xlsx  
...
```

Each file contains detailed step-level logs of the agent–market interaction.

4) How to experiment with the model

There are **three main places** where you can modify the system.

A) Change RL / PPO hyperparameters

Edit the default values inside the `make_config()` function in `ML_main.py`.

Typical parameters you may want to adjust:

- PPO hyperparameters: `learning_rate`, `clip_range`, `ent_coef`, `gamma`, `gae_lambda`,
`n_steps`, `batch_size`, `n_epochs`
- Training horizon: `train_episodes_chunks`
- Episode structure: `episode_len_t`, `window_size`
- Network architecture: `pi_arch`, `vf_arch`, `activation_fn`

Multiple experiments are defined via a list of configs:

```
runs = [  
    make_config(train_episodes_chunks=200),  
    make_config(train_episodes_chunks=300),  
    make_config(train_episodes_chunks=600),  
  
    make_config(learning_rate=3e-4),  
    make_config(ent_coef=0.005),  
    make_config(clip_range=0.2),  
]
```

To add a new experiment, simply add another `make_config(...)` entry.

B) Change the market price process

Edit the function `build_market_matrices()` in `prices_process.py`.

This function defines the **data-generating process** of prices, including:

- Long-run mean log-prices
- Mean reversion matrix
- Instant price impact of trades
- Persistent (decaying) impact of past trades
- Cross-good delayed impact
- Shock covariance and correlation structure
- Market volume / depth parameters

Examples:

- Increase volatility → increase shock variances
- Faster mean reversion → reduce diagonal entries of the reversion matrix
- Stronger market impact → increase instant or persistent impact coefficients
- More cross-good interaction → adjust cross-impact matrices

This is the main place to experiment with different market environments.

C) Change policy architecture

Still inside `make_config()` in `ML_main.py`, you can modify:

- Number of hidden layers
- Size of each layer
- Activation function

Example:

```
pi_arch = (256, 256)
vf_arch = (256, 256)
activation_fn = torch.nn.Tanh
```

5) Analysis and plots

All post-training analysis is handled in `ML/Results_analysis.py`.

A) Register raw output files

Inside the function `import_files_raw()`, list all Excel files you want to analyze:

```

FILES = {
    "120k data": "model_value_lenght_120000_raw.xlsx",
    "60k data": "model_value_lenght_60000_raw.xlsx",
    "40k data": "model_value_lenght_40000_raw.xlsx",

    "higher learning rate":
    "model_value_lenght_60000_learning_rate_0.0003_raw.xlsx",
    "lower entropy coefficient":
    "model_value_lenght_60000_ent_coef_0.005_raw.xlsx",
    "higher clip range": "model_value_lenght_60000_clip_range_0.2_raw.xlsx",
}

```

The keys are the labels used in plots; the values are the Excel filenames.

B) Run analysis

From the project root:

```
python ML/Results_analysis.py
```

This script:

1. Loads all registered Excel files
 2. Builds normalized variables (wealth, cash)
 3. Aggregates data across time, episodes, and learning phases
 4. Saves a clean long-format dataset
 5. Generates plots into `ML/plots/`
-

6) Outputs

Training outputs

- Raw logs: `model_value_lenght_*_raw.xlsx`

Analysis outputs

- `CLEAN_ALL.csv` (cached aggregated dataset)
 - Multiple `.png` figures in `ML/plots/`, including:
 - Wealth and cash evolution
 - Reward and penalty diagnostics
 - Within-episode dynamics
 - Cross-model comparisons
-

7) Typical workflow

1. Decide what to test (training horizon, PPO hyperparameters, architecture, market structure)
 2. Modify:
 3. `make_config()` and the runs list in `ML_main.py`
 4. `build_market_matrices()` in `prices_process.py`
 5. Run training: `python ML/ML_main.py`
 6. Register new output files in `import_files_raw()`
 7. Run analysis: `python ML/Results_analysis.py`
 8. Inspect figures in `ML/plots/`
-

This setup is designed as a **controlled experimental testbed**: the market is fully synthetic, the agent is trained under fixed constraints, and all dynamics are logged for reproducible analysis.